

agg_watch_data_classifier

May 8, 2022

1 Aggregate Watch Data Classification Project

The goal of this project is to investigate and utilize the data collected from a personal smartwatch to provide daily workout recommendations. Using the data collected from the Withings brand watch, we want to predict whether or not a person will have a successful workout on a given day. Providing this insight to users in the morning could provide valuable information about how the user could structure their day or provide the necessary motivation to make a workout routine become a workout habit. The idea of a “successful workout” will be investigated as well as which data provides insights in workout performance during the next day.

As an initial analysis, the data that has been aggregated by day will be used to determine whether or not it is a good predictor of a workout the following day, additionally the sleep data will be organized and cleaned to provide additional insights

Convert to PDF: jupyter nbconvert --execute --to pdf agg_watch_data_classifier.ipynb

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
```

1.1 Importing the Data

In this case, we are only looking at the data that has already been aggregated by day and not every file included in the watch_data folder

```
[2]: distance = pd.read_csv("../watch_data/aggregates_distance.csv", header=0)

passive_calories = pd.read_csv("../watch_data/aggregates_calories_passive.csv",
    ↪header=0)
active_calories = pd.read_csv("../watch_data/aggregates_calories_earned.csv",
    ↪header=0)

steps = pd.read_csv("../watch_data/aggregates_steps.csv", header=0)

sleep_data = pd.read_csv("../watch_data/sleep.csv", header=0)

workouts = pd.read_csv("../watch_data/activities.csv", header=0)
```

```
# Aggregate if I want to do same operation on all DataFrames
datasets = [distance, passive_calories, active_calories, steps, sleep_data,
↳workouts]
```

1.2 Data Preparation/Cleaning

Now that we have imported all of the relevant data, we need to clean the data and prepare it for model fitting

```
[3]: def convert_to_datetime(df:pd.DataFrame):
    if 'date' in df.columns:
        df['date'] = pd.to_datetime(df['date'], infer_datetime_format=True)
    elif 'from' in df.columns:
        df['from'] = pd.to_datetime(df['from'], infer_datetime_format=True,
↳utc=True)
        df['to'] = pd.to_datetime(df['to'], infer_datetime_format=True,
↳utc=True)
    else:
        print('No columns defined as date/from/to in {}'.format(df))

# Start with the simple files, check for nans and turn the date columns into
↳datetime types
for df in datasets:
    convert_to_datetime(df)
```

```
[4]: distance.rename(columns={'value':'distance'}, inplace=True)
passive_calories.rename(columns={'value':'passive calories'}, inplace=True)
active_calories.rename(columns={'value':'active calories'}, inplace=True)
steps.rename(columns={'value':'steps'}, inplace=True)
```

1.2.1 Now, we can focus on the “workouts” dataset. This dataset will require more work to get into a useable format. We will first drop the columns that contain only NaN values that provide no additional information that can be inferred or

```
[5]: workouts.drop(['from (manual)', 'to (manual)', 'GPS', 'Modified'], axis=1,
↳inplace=True)
```

```
[6]: workouts.head()
```

```
[6]:
```

		from	to	Timezone \
0	2020-05-10 12:54:00+00:00	2020-05-10 15:12:00+00:00	America/New_York	
1	2020-05-22 01:31:24+00:00	2020-05-22 02:16:10+00:00	America/New_York	
2	2020-05-22 01:31:24+00:00	2020-05-22 02:16:10+00:00	America/New_York	
3	2020-05-22 20:35:31+00:00	2020-05-22 21:41:57+00:00	America/New_York	
4	2020-05-22 20:35:31+00:00	2020-05-22 21:41:57+00:00	America/New_York	

	Activity type	Data
0	Walking	{"calories":390.5407409668,"effduration":8280,...
1	Yoga	{"calories":321.34295654297,"effduration":1080...
2	Other	{}
3	Other	{"calories":85.029991149902,"effduration":3300...
4	Other	{}

1.2.2 It's clear that a lot of information is located in the “Data” column of the workouts DataFrame and needs to be unpacked. Many of the elements in the Data column are the empty set, we also need to investigate this.

```
[7]: # Get the total number of times the Data column is the empty array
(workouts['Data'] == '{}').sum()
```

```
[7]: 241
```

```
[8]: # Maybe the "Other" workout types result in this output, however there are only
      ↳ 114 "Other" workouts registered
workouts['Activity type'].value_counts()
```

```
[8]: Walking          442
     Gym class       304
     Other           114
     Yoga             57
     Cycling          10
     Indoor Cycling    6
     Weights           6
     Tennis            4
     Windsurfing        1
     Elliptical         1
     Name: Activity type, dtype: int64
```

```
[9]: # Showing which type of workouts results in the null bracket for the workout
workouts[workouts['Data'] == '{}']['Activity type'].value_counts()
```

```
[9]: Gym class       146
     Other           75
     Yoga            14
     Cycling          6
     Name: Activity type, dtype: int64
```

1.2.3 The empty bracket categories are somewhat random, let's unpack the Data column to see what it contains for the different workouts

```
[10]: import ast
      ast.literal_eval(workouts['Data'].iloc[0])
```

```
[10]: {'calories': 390.5407409668,
      'effduration': 8280,
      'intensity': 0,
      'manual_distance': -1,
      'manual_calories': -1,
      'pause_duration': 0,
      'steps': 12591,
      'distance': 9793.205078125,
      'metcumul': 431.79629516602}
```

```
[11]: # Here we see that the 'effduration' category is simply the number of seconds,
      ↪ the workout is
      (workouts['to'].iloc[0] - workouts['from'].iloc[0]).total_seconds()
```

```
[11]: 8280.0
```

1.2.4 As a first investigation before diving into the 'Data' column too heavily as it changes for different workout types, we can simply compute the workout duration and save it as a new column

```
[12]: # Create the column that generates a TimeDelta datetime object
      workouts['Duration'] = (workouts['to'] - workouts['from']) / np.timedelta64(1,
      ↪ 's')
      workouts['Duration']
```

```
[12]: 0      8280.0
      1      2686.0
      2      2686.0
      3      3986.0
      4      3986.0
      ...
      940     2815.0
      941     1046.0
      942     1046.0
      943     2393.0
      944     2393.0
      Name: Duration, Length: 945, dtype: float64
```

1.2.5 Here it becomes more clear that the workouts with the empty brackets are simply duplicates of the previous workouts without the data, so we can remove each of these rows from the dataset

```
[13]: workouts = workouts[workouts['Data'] != '{}']
```

1.2.6 Now let's make a few changes to simplify the workouts data.

1. Replace the “from” and “to” columns with one date that indicates the date of the workout in addition to the “Duration” column
2. Remove the “Data” column (further investigation at a later point)

Care is needed because there may be multiple workouts in one day. In this case, if multiple workouts occur on the same day, they will be summed into one day and one duration value. In this case, we will also drop the activity type and timezone and simply find the total duration for each day that a workout was completed

```
[14]: basicworkouts = workouts.copy()
basicworkouts['date'] = pd.to_datetime(workouts['from'].dt.date)
basicworkouts.drop(['from', 'to', 'Data', 'Timezone', 'Activity type'], axis=1,
    ↪inplace=True)
basicworkouts.set_index('date', inplace=True)
```

```
[15]: # Sum workout duration on the days when there are more than one workout and
    ↪then consider a workout effective if it lasts longer than the 30 minutes
    ↪recommended by CDC
basicworkouts = basicworkouts.groupby(['date']).sum()
basicworkouts['Effective Workout'] = basicworkouts['Duration'] > 1800
basicworkouts['Effective Workout'].value_counts()
```

```
[15]: True      323
False     105
Name: Effective Workout, dtype: int64
```

1.3 Next, we need to clean that sleep_data file

```
[16]: sleep_data.head()
```

```
[16]:
```

		from	to	light (s)	deep (s)	\
0	2020-05-22	04:16:00+00:00	2020-05-22 12:19:00+00:00	17640	10260	
1	2020-05-23	06:47:00+00:00	2020-05-23 12:30:00+00:00	11880	8280	
2	2020-05-24	05:06:00+00:00	2020-05-24 12:13:00+00:00	11880	12900	
3	2020-05-25	04:57:00+00:00	2020-05-25 14:02:00+00:00	16620	15000	
4	2020-05-26	04:37:00+00:00	2020-05-26 12:23:00+00:00	14280	13560	

	rem (s)	awake (s)	wake up	Duration to sleep (s)	\
0	0	1080	3	120	
1	0	420	0	120	

2	0	840	2	120
3	0	1080	2	120
4	0	120	0	120

	Duration to wake up (s)	Snoring (s)	Snoring episodes	Average heart rate \
0	180	0	0	37
1	300	0	0	46
2	60	0	0	40
3	60	0	0	40
4	0	0	0	40

	Heart rate (min)	Heart rate (max)	Night events
0	33	55	NaN
1	36	95	NaN
2	32	73	NaN
3	32	77	NaN
4	31	77	NaN

```
[17]: sleep_data.describe()
```

```
[17]:
```

	light (s)	deep (s)	rem (s)	awake (s)	wake up \
count	564.000000	564.000000	564.0	564.000000	564.000000
mean	13637.854610	12890.441489	0.0	1341.386525	1.765957
std	3670.772948	3078.419841	0.0	872.533436	1.421171
min	0.000000	0.000000	0.0	0.000000	0.000000
25%	11565.000000	10980.000000	0.0	720.000000	1.000000
50%	14010.000000	13020.000000	0.0	1200.000000	1.000000
75%	16020.000000	15060.000000	0.0	1800.000000	2.000000
max	23580.000000	21180.000000	0.0	4200.000000	8.000000

	Duration to sleep (s)	Duration to wake up (s)	Snoring (s) \
count	564.000000	564.000000	564.0
mean	143.406028	187.553191	0.0
std	134.952251	253.224405	0.0
min	0.000000	0.000000	0.0
25%	120.000000	0.000000	0.0
50%	120.000000	120.000000	0.0
75%	120.000000	300.000000	0.0
max	1740.000000	1860.000000	0.0

	Snoring episodes	Average heart rate	Heart rate (min) \
count	564.0	564.000000	564.000000
mean	0.0	40.670213	32.686170
std	0.0	7.675555	6.381235
min	0.0	0.000000	0.000000
25%	0.0	38.000000	32.000000
50%	0.0	40.000000	32.000000

75%	0.0	43.000000	33.000000
max	0.0	112.000000	109.000000

	Heart rate (max)	Night events
count	564.000000	0.0
mean	78.333333	NaN
std	17.838795	NaN
min	0.000000	NaN
25%	70.000000	NaN
50%	80.000000	NaN
75%	88.000000	NaN
max	137.000000	NaN

1.3.1 It's clear from the sleep data that REM sleep is not recorded, as well as snoring, snoring episodes, and night events. These columns can be dropped. Additionally, "from" and "to" columns are not necessary as we simply want the date that the sleep occurred that corresponds with a workout later that day

```
[18]: sleep_data.drop(['rem (s)', 'Snoring (s)', 'Snoring episodes', 'Night events'],
    ↪axis=1, inplace=True)
```

```
[19]: sleep_data['date'] = pd.to_datetime(sleep_data["from"].dt.date)
```

```
[20]: # If we look at the values for the dates, we see that we have one day in which
    ↪there are 2 recorded sleeps! Let's investigate this date
    sleep_data['date'].value_counts()
```

```
[20]: 2022-01-27    2
    2021-04-03    1
    2021-11-11    1
    2021-08-26    1
    2021-05-20    1
    ..
    2021-03-04    1
    2021-05-05    1
    2021-10-27    1
    2021-09-06    1
    2022-02-28    1
    Name: date, Length: 563, dtype: int64
```

```
[21]: sleep_data[sleep_data['date'] == '20220127']
```

```
[21]:
```

		from	to	light (s)	deep (s)	\
498	2022-01-27	01:52:00+00:00	2022-01-27 10:49:00+00:00	14520	13860	
499	2022-01-27	15:47:00+00:00	2022-01-27 21:28:00+00:00	2940	16620	

	awake (s)	wake up	Duration to sleep (s)	Duration to wake up (s)	\
--	-----------	---------	-----------------------	-------------------------	---

498	3840	4	120	300
499	900	1	120	600

	Average heart rate	Heart rate (min)	Heart rate (max)	date
498	40	32	80	2022-01-27
499	112	109	115	2022-01-27

1.3.2 After investigating, it's clear that the duplicate was due to a long nap I took on vacation :) I will remove this sleep record from the data set in this case. In other applications, it may become necessary to create some outlier detection to determine when a sleep event occurs outside of the usual time.

```
[22]: sleep_data.drop(499, inplace=True)
sleep_data.drop(['from', 'to'], axis=1, inplace=True)
sleep_data.set_index('date', inplace=True)
```

```
[23]: # Cleaned sleep data
sleep_data
```

```
[23]:
```

	light (s)	deep (s)	awake (s)	wake up	Duration to sleep (s)	\
date						
2020-05-22	17640	10260	1080	3	120	
2020-05-23	11880	8280	420	0	120	
2020-05-24	11880	12900	840	2	120	
2020-05-25	16620	15000	1080	2	120	
2020-05-26	14280	13560	120	0	120	
...	
2022-04-13	13500	14040	1560	2	120	
2022-04-14	18720	14640	120	0	120	
2022-04-15	16320	15120	1980	2	120	
2022-04-16	16800	15060	1740	2	120	
2022-04-17	15600	10500	180	0	120	

	Duration to wake up (s)	Average heart rate	Heart rate (min)	\
date				
2020-05-22	180	37	33	
2020-05-23	300	46	36	
2020-05-24	60	40	32	
2020-05-25	60	40	32	
2020-05-26	0	40	31	
...	
2022-04-13	0	41	32	
2022-04-14	0	44	33	
2022-04-15	0	38	33	
2022-04-16	0	36	32	
2022-04-17	60	39	33	

	Heart rate (max)
date	
2020-05-22	55
2020-05-23	95
2020-05-24	73
2020-05-25	77
2020-05-26	77
...	...
2022-04-13	114
2022-04-14	84
2022-04-15	80
2022-04-16	53
2022-04-17	73

[563 rows x 9 columns]

1.4 Now that the more complicated datasets have been cleaned, let's quickly clean the distance, passive_calories, active_calories, and steps sets

```
[24]: distance.set_index('date', inplace=True)
passive_calories.set_index('date', inplace=True)
active_calories.set_index('date', inplace=True)
steps.set_index('date', inplace=True)
```

```
[25]: distance.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 711 entries, 2022-04-17 to 2020-05-07
Data columns (total 1 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   distance    711 non-null    float64
dtypes: float64(1)
memory usage: 11.1 KB
```

```
[26]: basicworkouts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 428 entries, 2020-05-10 to 2022-04-15
Data columns (total 2 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Duration            428 non-null    float64
1   Effective Workout    428 non-null    bool
dtypes: bool(1), float64(1)
memory usage: 7.1 KB
```

1.5 Now, let's join the datasets together on the date

```
[27]: X_data = pd.merge(distance, passive_calories, how='inner', on='date')
X_data = pd.merge(X_data, active_calories, how='inner', on='date')
X_data = pd.merge(X_data, steps, how='inner', on='date')
X_data = pd.merge(X_data, sleep_data, how='inner', on='date')
X_data = pd.merge(X_data, basicworkouts, how='inner', on='date')
X_data
```

```
[27]:
```

	distance	passive calories	active calories	steps	light (s)	\
date						
2022-04-15	2579.809	1869.678	106.592	3082	16320	
2022-04-14	5363.195	1869.678	208.608	6596	18720	
2022-04-13	2611.168	1869.678	102.255	3212	13500	
2022-04-12	2773.922	1869.678	112.184	3429	14400	
2022-04-10	3565.152	1869.678	139.220	4370	12540	
...	
2020-05-28	5218.729	1914.158	206.033	6483	16020	
2020-05-27	2222.013	1914.158	87.538	2596	15120	
2020-05-25	1772.393	1914.158	69.988	2052	16620	
2020-05-24	12708.024	1914.158	502.375	16357	11880	
2020-05-22	5737.927	1914.158	226.621	6758	17640	

	deep (s)	awake (s)	wake up	Duration to sleep (s)	\
date					
2022-04-15	15120	1980	2	120	
2022-04-14	14640	120	0	120	
2022-04-13	14040	1560	2	120	
2022-04-12	12420	1500	1	180	
2022-04-10	13440	2520	3	120	
...	
2020-05-28	13740	1980	1	120	
2020-05-27	8820	1260	3	120	
2020-05-25	15000	1080	2	120	
2020-05-24	12900	840	2	120	
2020-05-22	10260	1080	3	120	

	Duration to wake up (s)	Average heart rate	Heart rate (min)	\
date				
2022-04-15	0	38	33	
2022-04-14	0	44	33	
2022-04-13	0	41	32	
2022-04-12	480	36	33	
2022-04-10	60	38	33	
...	
2020-05-28	1680	41	33	
2020-05-27	240	38	33	

2020-05-25	60	40	32
2020-05-24	60	40	32
2020-05-22	180	37	33

date	Heart rate (max)	Duration	Effective Workout
2022-04-15	80	3439.0	True
2022-04-14	84	2815.0	True
2022-04-13	114	3574.0	True
2022-04-12	47	1080.0	False
2022-04-10	69	6548.0	True
...
2020-05-28	99	1980.0	True
2020-05-27	47	2647.0	True
2020-05-25	77	1560.0	False
2020-05-24	73	8676.0	True
2020-05-22	55	7932.0	True

[350 rows x 15 columns]


2. Exploratory Data Analysis

2.0.1 Now that the data has been cleaned to a usable format, we can briefly explore the data before applying different ML techniques. It should be noted that this data will only include the days in which data exists for each of the previous datasets, i.e. days in which I had a workout, recorded my sleep, and other data exists (recorded automatically every day)

```
[28]: fig, axes = plt.subplots(figsize=(12, 6))
      axes.scatter(X_data.index, X_data['steps'])
      axes.set_xlabel('Date')
      axes.set_ylabel('Steps')

      axes.set_title("My Steps on Workout Days Over Time")
```

```
[28]: Text(0.5, 1.0, 'My Steps on Workout Days Over Time')
```



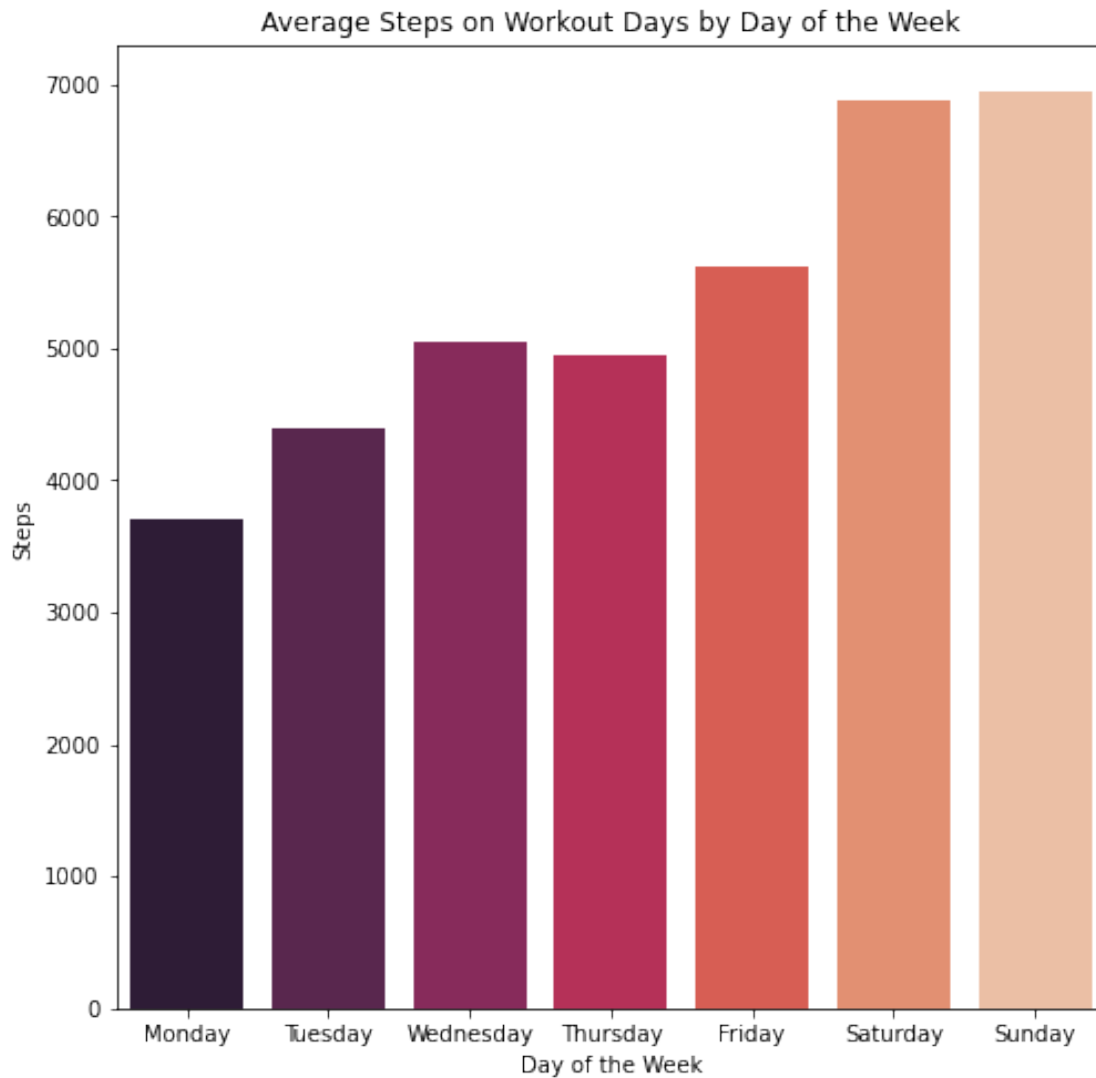
```
[29]: cats = [ 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
        steps_by_day = X_data.groupby(X_data.index.day_name()).mean().reindex(cats)
```

```
[30]: import seaborn as sb

fig, axes = plt.subplots(figsize=(8, 8))
sb.barplot(x = steps_by_day.index, y = steps_by_day['steps'], palette="rocket")
axes.set_xlabel('Day of the Week')
axes.set_ylabel('Steps')

axes.set_title("Average Steps on Workout Days by Day of the Week")
```

```
[30]: Text(0.5, 1.0, 'Average Steps on Workout Days by Day of the Week')
```



```
[31]: X_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 350 entries, 2022-04-15 to 2020-05-22
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   distance             350 non-null    float64
1   passive calories     350 non-null    float64
2   active calories      350 non-null    float64
3   steps                350 non-null    int64
4   light (s)            350 non-null    int64
5   deep (s)             350 non-null    int64
6   awake (s)            350 non-null    int64
```

```

7  wake up                350 non-null    int64
8  Duration to sleep (s)  350 non-null    int64
9  Duration to wake up (s) 350 non-null    int64
10 Average heart rate     350 non-null    int64
11 Heart rate (min)       350 non-null    int64
12 Heart rate (max)       350 non-null    int64
13 Duration               350 non-null    float64
14 Effective Workout      350 non-null    bool
dtypes: bool(1), float64(4), int64(10)
memory usage: 41.4 KB


```

```
[32]: X_data.hist(bins=50, figsize=(20,15))
```

```

[32]: array([[<AxesSubplot:title={'center':'distance'}>,
             <AxesSubplot:title={'center':'passive calories'}>,
             <AxesSubplot:title={'center':'active calories'}>,
             <AxesSubplot:title={'center':'steps'}>],
           [<AxesSubplot:title={'center':'light (s)'}>,
             <AxesSubplot:title={'center':'deep (s)'}>,
             <AxesSubplot:title={'center':'awake (s)'}>,
             <AxesSubplot:title={'center':'wake up'}>],
           [<AxesSubplot:title={'center':'Duration to sleep (s)'}>,
             <AxesSubplot:title={'center':'Duration to wake up (s)'}>,
             <AxesSubplot:title={'center':'Average heart rate'}>,
             <AxesSubplot:title={'center':'Heart rate (min)'}>],
           [<AxesSubplot:title={'center':'Heart rate (max)'}>,
             <AxesSubplot:title={'center':'Duration'}>, <AxesSubplot:>,
             <AxesSubplot:>]], dtype=object)

```



3. Feature Exploration/Engineering

3.1 Many of the techniques have been chosen from the text “Hands on Machine Learning”

```
[33]: corr_mat = X_data.corr()
      corr_mat['Effective Workout'].sort_values(ascending=False)
```

```
[33]: Effective Workout      1.000000
      Duration              0.645952
      light (s)             0.124674
      Duration to wake up (s) 0.108354
      distance              0.092069
      active calories        0.090891
      steps                 0.090876
      deep (s)              0.034190
      awake (s)             0.032055
      wake up               0.017298
      Heart rate (max)       0.005117
      passive calories       -0.000140
```

```

Duration to sleep (s)      -0.016959
Average heart rate        -0.033412
Heart rate (min)          -0.051027
Name: Effective Workout, dtype: float64

```

3.2 After looking at the correlations, it's clear that some of the data would not make sense to predict whether or not someone will have an effective workout. After consideration, it makes sense to predict whether an effective workout will occur or not based off of information from the sleep data OR data from the previous day. Because of this, the following steps will be made to modify the data:

1. The Duration, distance, active calories, steps, and passive calories categories will be removed from the current day as they occur concurrently with the current day's workout and may be confounding variables
2. The distance, active/passive calories, and steps from the previous day will be added in as possible influence over an effective workout or not
3. The total time asleep is added as a feature

```

[34]: # Remove confounding variables
x_test = X_data.drop(['Duration', 'distance', 'active calories', 'steps',
↳ 'passive calories'], axis=1)

```

```

[35]: # Add total time asleep as a column
x_test['total sleep'] = X_data['light (s)'] + X_data['deep (s)']

```

```

[36]: from datetime import datetime, timedelta
prev_day_index = x_test.index - timedelta(days=1)
x_test['prevday steps'] = steps['steps'].loc[prev_day_index].values
x_test['prevday active cals'] = active_calories['active calories'].
↳ loc[prev_day_index].values
x_test['prevday passive cals'] = passive_calories['passive calories'].
↳ loc[prev_day_index].values
x_test['prevday distance'] = distance['distance'].loc[prev_day_index].values

```

```

[37]: corr_mat = x_test.corr()
corr_mat['Effective Workout'].sort_values(ascending=False)

```

```

[37]: Effective Workout      1.000000
light (s)                   0.124674
total sleep                  0.111419
Duration to wake up (s)     0.108354
deep (s)                    0.034190
awake (s)                   0.032055
wake up                     0.017298
Heart rate (max)            0.005117
Duration to sleep (s)       -0.016959

```



```

prevday passive cals      -0.023298
Average heart rate        -0.033412
Heart rate (min)          -0.051027
prevday distance          -0.154032
prevday steps             -0.155316
prevday active cals       -0.156304
Name: Effective Workout, dtype: float64

```

3.2.1 Now, we can plot the correlation matrix with all of our data we will use to predict an effective workout.

From this data, the main correlation to “Effective Workout” come from the sleep data with the highest correlation related to light sleep.

```
[38]: corr_mat.style.background_gradient(cmap='coolwarm')
```

```
[38]: <pandas.io.formats.style.Styler at 0x7fc90031f670>
```

```
[39]: x_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 350 entries, 2022-04-15 to 2020-05-22
Data columns (total 15 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   light (s)                            350 non-null    int64
 1   deep (s)                             350 non-null    int64
 2   awake (s)                            350 non-null    int64
 3   wake up                              350 non-null    int64
 4   Duration to sleep (s)                350 non-null    int64
 5   Duration to wake up (s)              350 non-null    int64
 6   Average heart rate                    350 non-null    int64
 7   Heart rate (min)                     350 non-null    int64
 8   Heart rate (max)                     350 non-null    int64
 9   Effective Workout                     350 non-null    bool
10   total sleep                           350 non-null    int64
11   prevday steps                         350 non-null    int64
12   prevday active cals                   350 non-null    float64
13   prevday passive cals                  350 non-null    float64
14   prevday distance                      350 non-null    float64
dtypes: bool(1), float64(3), int64(11)
memory usage: 41.4 KB

```

3.3 Prepping Data for Fitting Classifiers

```
[40]: # Getting a test set
from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(x_test, test_size=0.2, random_state=42)
```

```
[41]: y_train = train_set['Effective Workout'].values.astype(int)
x_train = train_set.drop(['Effective Workout'], axis=1)

y_test = test_set['Effective Workout'].values.astype(int)
x_test = test_set.drop(['Effective Workout'], axis=1)
```

```
[42]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

num_pipeline = Pipeline([
    ('std_scaler', StandardScaler()),
])
x_train_prepared = num_pipeline.fit_transform(x_train)
x_test_prepared = num_pipeline.fit_transform(x_test)
```

3.4 Stochastic Gradient Descent classifier

```
[43]: # Training the classifier in the standard way
from sklearn.linear_model import SGDClassifier

sgd_clf = SGDClassifier(random_state=42)
sgd_clf.fit(x_train_prepared, y_train)
```

```
[43]: SGDClassifier(random_state=42)
```

```
[44]: # Getting the cross-validation score
from sklearn.model_selection import cross_val_score
cross_val_score(sgd_clf, x_train_prepared, y_train, cv=3, scoring="accuracy")
```

```
[44]: array([0.61702128, 0.65591398, 0.65591398])
```

```
[45]: # Getting the confusion matrix
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix

y_train_pred = cross_val_predict(sgd_clf, x_train_prepared, y_train, cv=3)

y_scores = cross_val_predict(sgd_clf, x_train_prepared, y_train, cv=3,
    ↪method="decision_function")
```

```
[46]: confusion_matrix(y_train, y_train_pred)
```

```
[46]: array([[ 22,  49],
          [ 51, 158]])
```

3.4.1 Looking at the results we see:

1. Our classifier performs equally poorly at Type 1 and 2 errors and that our resulting precision and recall is very similar. This may mean that our SGD classifier can do no better than what is shown here without tradeoffs between precision and recall
2. The average cross-validation accuracy of around 0.63 is somewhat better than random chance, but doesn't tell the full story because the precision and recall are decent compared with the accuracy
3. Precision: 0.763
4. Recall : 0.756
5. F1 score : 0.7596
6. The classifier is not good at predicting true negative, or cases when a poor workout is expected. This could mean that we need to train a different model, or that the input features are simply not good predictors of a good workout or not

```
[47]: from sklearn.metrics import precision_score, recall_score, f1_score

precision_score(y_train, y_train_pred)
```

```
[47]: 0.7632850241545893
```

```
[48]: recall_score(y_train, y_train_pred)
```

```
[48]: 0.7559808612440191
```

```
[49]: f1_score(y_train, y_train_pred)
```

```
[49]: 0.7596153846153845
```

```
[50]: # What would the accuracy be if we simply guessed a good workout every time?
sum(np.ones(len(y_train)) == y_train)/len(y_train)
```

```
[50]: 0.7464285714285714
```

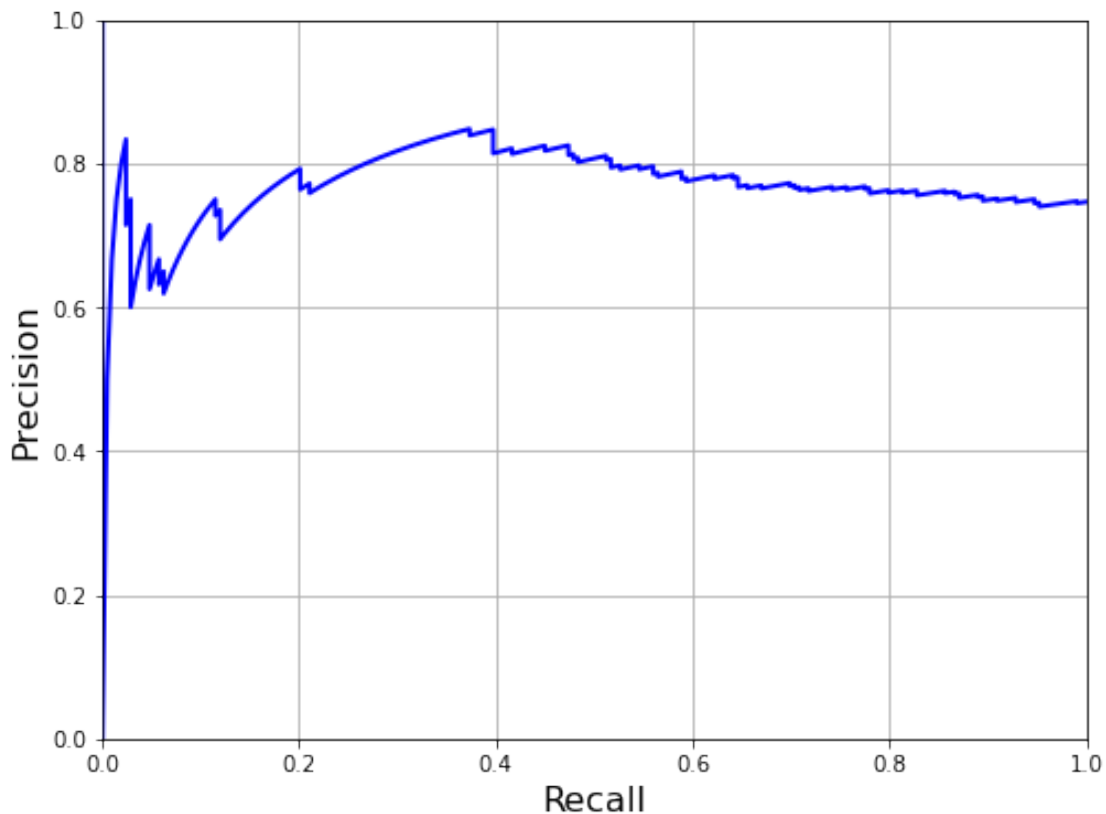
```
[51]: from sklearn.metrics import precision_recall_curve

precisions, recalls, thresholds = precision_recall_curve(y_train, y_scores)

def plot_precision_vs_recall(precisions, recalls):
    plt.plot(recalls, precisions, "b-", linewidth=2)
    plt.xlabel("Recall", fontsize=16)
    plt.ylabel("Precision", fontsize=16)
```

```
plt.axis([0, 1, 0, 1])
plt.grid(True)

plt.figure(figsize=(8, 6))
plot_precision_vs_recall(precisions, recalls)
plt.show()
```



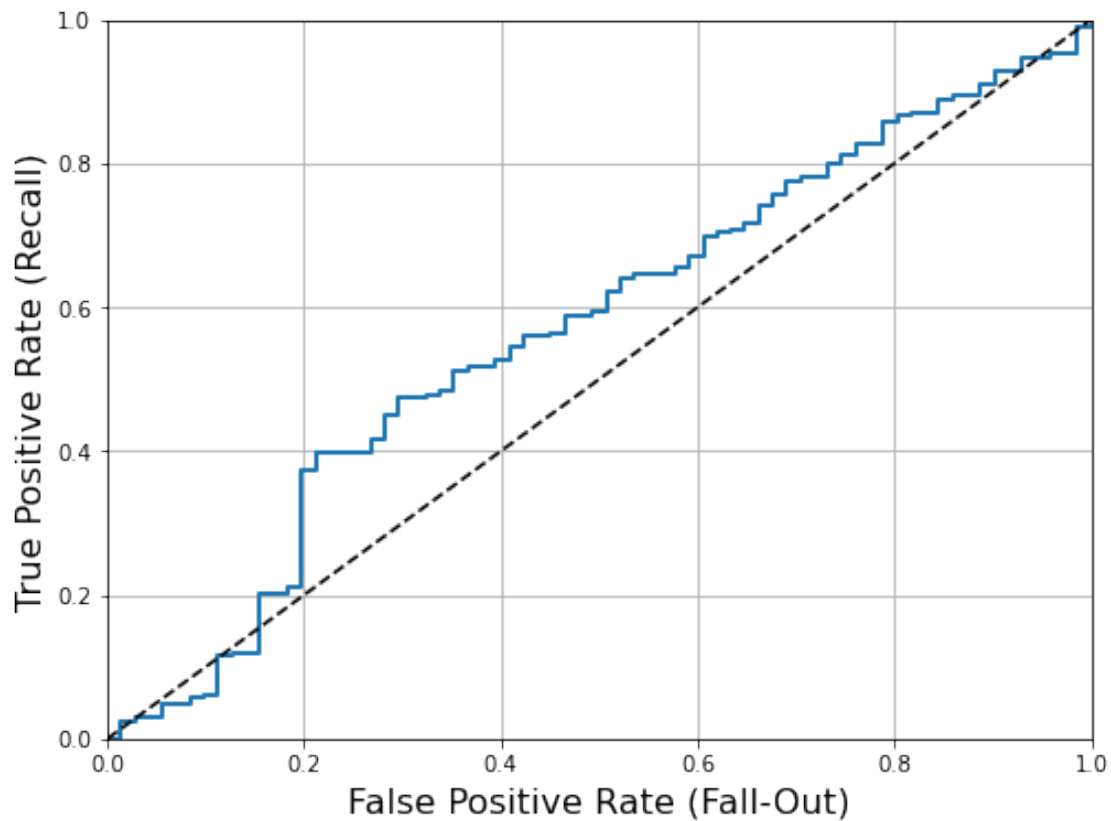
3.4.2 We can look at the ROC curve to see the performance of our SGD classifier

```
[52]: from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_train, y_scores)

def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'k--')
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate (Fall-Out)', fontsize=16)
    plt.ylabel('True Positive Rate (Recall)', fontsize=16)
```

```
plt.grid(True)

plt.figure(figsize=(8, 6))
plot_roc_curve(fpr, tpr)
plt.show()
```



```
[53]: from sklearn.metrics import roc_auc_score
roc_auc_score(y_train, y_scores)
```

[53]: 0.5665476110250017

3.5 From this initial investigation, the current classifier performs poorly. This could be due to a number of factors:

1. Small datasets (more data may differentiate bad workouts from good workouts more)
2. Class imbalance (75% are effective workouts and 25% are not)
3. Incorrect metrics (perhaps a 30 minute workout is not the sure-fire metric that was expected)

3.6 First, other binary classifiers will be tested and then further analysis will be done